

STATS300A - Lecture 9

Dominik Rothenhaeusler
Scribed by Michael Howes

10/18/21

Contents

| | |
|---|----------|
| 1 Overview | 1 |
| 2 Stengths of the Bayesian approach | 1 |
| 2.1 Bayesian recursion | 2 |
| 2.2 Hierarchical models and empirical Bayes | 2 |
| 2.3 Why Bayes estimators? | 3 |
| 3 Minimax estimators | 4 |

1 Overview

We have been studying optimal estimation. We have tried multiple things:

- (a) Uniform comparisons.
- (b) Restricting the class of estimators.
- (c) Collapsing the risk.
 - i. Bayesian estimators.
 - ii. Minimax estimators.

Today we will discuss:

- Some strengths of Bayesian techniques.
- Minimax risk estimation.

2 Stengths of the Bayesian approach

Lemma 1. [TPE 4.14] *Let Q be the marginal distribution of X . That is*

$$Q(A) = \int_{\Omega} \mathbb{P}_{\theta}(X \in A) d\Lambda.$$

If the loss is strictly convex in d , then the Bayes estimator δ_{Λ} is unique a.s. \mathbb{P}_{θ} for all $\theta \in \Omega$ if

- (a) $r(\Lambda, d_{\Lambda}) < \infty$.
- (b) *If $A \subseteq \mathcal{X}$ and $Q(A) = 0$, then $\mathbb{P}_{\theta}(A) = 0$ for all $\theta \in \Omega$.*

See the textbook for a proof. Note that if the following all hold, then we can conclude that (b) holds above.

- (a) Ω is an open subset of \mathbb{R}^k .
- (b) The map $\theta \rightarrow \mathbb{P}_\theta(A)$ is continuous for all A .
- (c) $\pi(\theta) > 0$ for all $\theta \in \Omega$.
- (d) $\theta \rightarrow \pi(\theta)$ is continuous.

2.1 Bayesian recursion

Suppose $\theta \sim \Lambda$ and $X_i \stackrel{\text{iid}}{\sim} \mathbb{P}_\theta$. We can then update the prior sequentially. Note that the posterior for $m < n$ observations is

$$p(\theta|X_1, \dots, X_n) \propto \text{likelihood} \times \text{prior} = p(X_1, \dots, X_m|\theta)\pi(\theta).$$

The posterior for the full sample of n observations is

$$\begin{aligned} p(\theta|X_1, \dots, X_n) &\propto \text{likelihood} \times \text{prior} \\ &= p(X_1, \dots, X_n|\theta)\pi(\theta) \\ &= p(X_1, \dots, X_m|\theta)p(X_{m+1}, \dots, X_n|\theta)\pi(\theta) \\ &\propto p(X_{m+1}, \dots, X_n|\theta)p(\theta|X_1, \dots, X_m). \end{aligned}$$

Thus we can think of X_{m+1}, \dots, X_n as new data and the posterior $p(\theta|X_1, \dots, X_m)$ as a new prior. It follows that we can compute posteriors recursively by changing the prior. This allows for speedy calculations.

Example 1. Suppose $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta, \sigma^2)$ where σ^2 is known and $\theta \sim \mathcal{N}(\mu, b^2)$. We saw that

$$\text{posterior} \propto \exp \left\{ -\frac{1}{2}\theta^2 w_n + \theta \bar{w}_n \right\},$$

where

$$\begin{aligned} w_n &= \frac{n}{\sigma^2} + \frac{1}{b^2}, \\ \bar{w}_n &= \frac{\sum_{i=1}^n x_i}{\sigma^2} + \frac{\mu}{b^2}. \end{aligned}$$

The weights w_n satisfy the recursion

$$\begin{aligned} w_n &= \frac{1}{\sigma^2} + w_{n-1}, \\ \bar{w}_n &= \frac{x_n}{\sigma^2} + \bar{w}_{n-1}. \end{aligned}$$

Thus our updates are quick linear calculations. This has many applications.

2.2 Hierarchical models and empirical Bayes

We can use Bayesian ideas to model problems with repeat structure and pool information across observations. Suppose we have $\theta_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau^2)$ for $i = 1, \dots, p$ where $\tau^2 > 0$ is known (for now).

Suppose that we also have $X_i \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_i, 1)$. For example θ_i might be the effect of an experiment and X_i is the measured effect and our measurement errors are i.i.d. $\mathcal{N}(0, 1)$. Consider the loss

$$L(\theta, \delta) = \sum_{i=1}^p (\theta_i - \delta_i(X))^2.$$

One can show (see Keener section 11) that the Bayes estimator is given by

$$\delta_i(X) = \left(1 - \frac{1}{1 + \tau^2}\right) X_i = \frac{\tau^2}{1 + \tau^2} X_i.$$

That is we shrink our data towards 0. The amount of shrinkage depends on τ^2 . If τ^2 is small we shrink a lot, if it is large we shrink less. We will not prove that δ is the Bayes estimator but we will show that it has the lowest average risk of estimators of the form $\hat{\theta}_\alpha = \alpha X$ when $p = 1$. The risk of such an estimator is

$$\mathbb{E}[(\theta - \alpha X)^2] = \mathbb{E}[\theta^2] - 2\alpha\mathbb{E}[\theta X] + \alpha^2\mathbb{E}[X^2] = \tau^2 - 2\alpha\tau^2 + \alpha^2(\tau^2 + 1),$$

which is minimized when $\alpha(\tau^2 + 1) = \tau^2$, that is $\alpha = \frac{\tau^2}{1 + \tau^2}$. The shrinkage term is the ratio of the experiment noise and the measurement noise.

What if we don't know τ ? We can still estimate it. This is the idea behind empirical Bayes. We can estimate the variance of X_i since

$$X_i \stackrel{\text{iid}}{\sim} N(0, 1 + \tau^2),$$

Thus $\frac{1}{p} \sum_{j=1}^p X_j^2$ is unbiased for $1 + \tau^2$. This gives us the new estimator

$$\delta'_i(X) = \left(1 - \frac{p}{\sum_{j=1}^n X_j^2}\right) X_i.$$

This is an example of empirical Bayes when the prior is estimated from the data. We will revisit this estimator in the frequentist setting later in the course.

2.3 Why Bayes estimators?

- (a) Every admissible estimator is a Bayes estimator or a limit of Bayes estimators. That is for a sequence of priors Λ_n , $\delta_{\Lambda_n}(x) \rightarrow \delta(x)$ a.e. \mathbb{P}_θ .
- (b) Can incorporate prior experience and beliefs.
- (c) Quantification of uncertainty are sometimes easier to interpret.
- (d) Encode complex data structures.

How do we choose priors?

- (a) Subjective: previous knowledge.
- (b) Objective: Select an “uninformative prior” such as a Jeffery’s prior or a uniform prior.
- (c) Empirical: Estimate prior parameters from data.
- (d) Computational: Use conjugate priors.

As $n \rightarrow \infty$, the posterior is independent of the prior - Bernstein-von-Mises.

3 Minimax estimators

Given $X \sim \mathbb{P}_\theta$, θ fixed and unknown. Our goal is to find an estimator δ that minimizes

$$\sup_{\theta \in \Omega} R(\theta, \delta).$$

Such a δ is called a minimax estimator. How do we find such estimators? Note that

$$\sup_{\theta} R(\theta, \delta) \geq \int_{\Omega} R(\theta, \delta) d\Lambda,$$

for all probability distributions Λ . Thus our goal is to find the “worst prior.”

Definition 1. A prior Λ is called *least favourable* if $r_{\Lambda} \geq r_{\Lambda'}$ for any other prior Λ' .

Recall that $r_{\Lambda} = r(\Lambda, \delta_{\Lambda})$ where δ_{Λ} is the Bayes estimator for Λ .

Theorem 1. [TPE 5.1.4] Suppose δ_{Λ} is a Bayes estimator with $r_{\Lambda} = \sup_{\theta} R(\theta, \delta_{\Lambda})$, then

- (a) δ_{Λ} is minimax.
- (b) Λ is least favourable.
- (c) If δ_{Λ} is a unique Bayes estimator (a.e \mathbb{P}_{θ} for all $\theta \in \Omega$), then δ_{Λ} is the unique minimax estimator.

Proof. For (a) let δ be another estimator, then

$$\begin{aligned} \sup_{\theta} R(\theta, \delta) &\geq \int_{\Omega} R(\theta, \delta) d\Lambda \\ &\geq \int_{\Omega} R(\theta, \delta_{\Lambda}) d\Lambda \\ &= r_{\Lambda} \\ &= \sup_{\theta} R(\theta, \delta_{\Lambda}) \end{aligned}$$

Thus δ_{Λ} is minimax. For (b), let Λ' be another prior, then

$$\begin{aligned} r_{\Lambda'} &= \int_{\Omega} R(\theta, \delta_{\Lambda'}) d\Lambda' \\ &\leq \int_{\Omega} R(\theta, \delta_{\Lambda}) d\Lambda' \\ &\leq \sup_{\theta \in \Omega} R(\theta, \delta_{\Lambda}) \\ &= \int_{\Omega} R(\theta, \delta_{\Lambda}) d\Lambda \\ &= r_{\Lambda}. \end{aligned}$$

Thus Λ is least favourable. For (c), suppose that δ is a minimax estimator, then

$$\begin{aligned} r(\Lambda, \delta) &= \int_{\Omega} R(\theta, \delta) d\Lambda \\ &\leq \sup_{\theta \in \mathcal{L}} R(\theta, \delta) \\ &= \sup_{\theta \in \mathcal{L}} R(\theta, \delta_{\Lambda}) \\ &= r_{\Lambda}. \end{aligned}$$

Thus δ is the Bayes estimator of the prior Λ . It follows that if δ_{Λ} is the unique Bayes estimator for Λ , then δ_{Λ} is the unique minimax estimator. \square

Example 2. Suppose $X \sim \text{Bin}(n, \theta)$ and we want to find the minimax estimator of θ under squared error loss. Our goal is to find a Bayes estimator that has constant risk. This would automatically give

$$r_\Lambda = \sup_{\theta \in \Omega} R(\theta, \delta_\Lambda).$$

We previously say that if Θ has a $\text{Beta}(a, b)$ prior, then the Bayes estimator has the form

$$\delta_\Lambda(x) = \frac{x + a}{a + b + n}.$$

We wish to find a, b such that this estimator has constant risk. Note that

$$\begin{aligned} R(\theta, \delta_\Lambda) &= \mathbb{E}_\theta \left[\left(\frac{x + a}{n + a + b} - \theta \right)^2 \right] \\ &= \frac{1}{(n + a + b)^2} \mathbb{E}_\theta \left[(X + a - \theta(n + a + b))^2 \right] \\ &= \frac{1}{(n + a + b)^2} \mathbb{E}_\theta \left[(X - n\theta + a(1 - \theta) - b\theta)^2 \right] \\ &= \frac{1}{(n + a + b)^2} [\text{Var}_\theta(X - n\theta) + (a(1 - \theta) - b\theta)^2] \\ &= \frac{1}{(n + a + b)^2} [n\theta(1 - \theta) + (a(1 - \theta) - b\theta)^2]. \end{aligned}$$

Thus we wish to find a, b such that $n\theta(1 - \theta) + (a(1 - \theta) - b\theta)^2$ is constant in θ . The solution is $a = b = \frac{\sqrt{n}}{2}$. Thus

$$\delta_\Lambda(x) = \frac{x + \frac{\sqrt{n}}{2}}{n + \sqrt{n}},$$

is a minimax estimator for θ .