

STATS305A - Lecture 18

John Duchi
Scribed by Michael Howes

11/30/21

Contents

1	Announcements	1
2	M-estimation	1
2.1	Recap	1
2.2	Choosing the loss function	2
3	Outlier mitigation	2
4	Qunatile regression	4

1 Announcements

- HW4 out today. There will be one multi-part question and one optional question.
- Etude 4 out tonight. Lots of the material for etude 4 will be covered in class on Thursday.
- Both will be due on Thursday next week at 5pm.

2 M-estimation

2.1 Recap

To do regression with M-estimators we use losses other than the squared error to measure error and fit models. For some loss $l : \mathbb{R} \rightarrow \mathbb{R}_+$ we solve the minimization problem

$$\text{minimize } \frac{1}{n} \sum_{i=1}^n l(y_i - x_i^T b),$$

or sometimes

$$\text{minimize } \frac{1}{n} \sum_{i=1}^n l(y_i - x_i^T b) + \text{Reg}(b).$$

Where $\text{Reg}(b)$ is some sort of regularizer like the norm of b .

2.2 Choosing the loss function

In ordinary least squares we use the loss function $l(t) = \frac{1}{2}t^2$. There are two guiding principles we can use to choose alternative loss functions.

- (a) Suppose that we could solve $\min_f \mathbb{E}[l(Y - f(X))]$ so that $f(x) = \operatorname{argmin}_z \mathbb{E}[l(Y - z)|X = x]$. We would then like to choose the loss l so that the function f has a property that we care about. We will see this when we do quantile regression.
- (b) Another approach is to choose a loss function l so that we can avoid undue influence of outlying measurements/responses y . To do this we need a loss that is Lipschitz continuous.

Definition 1. A loss function l is *Lipschitz* if there exists $C \geq 0$ such that for all $t, s \in \mathbb{R}$,

$$|l(t) - l(s)| \leq C|t - s|.$$

The smallest such C is called the *Lipschitz constant* of l and is denoted by $\|l\|_{Lip}$.

Example 1. The absolute error loss $l(t) = |t|$ is Lipschitz with Lipschitz constant 1. We claim that $\operatorname{argmin}_t \mathbb{E}[|Y - t|] = \operatorname{med}(Y)$ where $\operatorname{med}(Y)$ denotes the median of Y . To see why this is true, let $R(t) = \mathbb{E}[|Y - t|]$. We will compute the left and right derivatives of R . Note that

$$l'_{\leftarrow}(t) = \text{right derivative of } l = \operatorname{sgn}_+(t) = \begin{cases} 1 & \text{if } t \geq 0, \\ -1 & \text{if } t < 0. \end{cases}$$

And

$$l'_{\rightarrow}(t) = \text{left derivative of } l = \operatorname{sgn}_-(t) = \begin{cases} 1 & \text{if } t > 0, \\ -1 & \text{if } t \leq 0. \end{cases}$$

Thus

$$R'_{\leftarrow}(t) = \mathbb{E}[\operatorname{sgn}_+(t - Y)] = \mathbb{P}(Y \leq t) - \mathbb{P}(Y > t),$$

and

$$R'_{\rightarrow}(t) = \mathbb{E}[\operatorname{sgn}_-(t - Y)] = \mathbb{P}(Y < t) - \mathbb{P}(Y \geq t).$$

Thus if $t < \operatorname{med}(Y)$, then $R'_{\rightarrow}(t) < 0$. If $t > \operatorname{med}(Y)$, then $R'_{\leftarrow}(t) > 0$ and thus the minimizer of $R(t)$ is $\operatorname{med}(Y)$.

3 Outlier mitigation

We will now talk more about guiding principle (b) which was about choosing losses that are robust against outliers. Define

$$L_n(b) := \frac{1}{n} \sum_{i=1}^n l(y_i - x_i^T b).$$

We want to know what happens to the minimizer of $L_n(b)$ when we replace (x_k, y_k) with (x_k^*, y_k^*) . Let

$$L_{n,k}(b) = \frac{1}{n} \sum_{i \neq k} l(y_i - x_i^T b) + \frac{1}{n} l(y_k^* - (x_k^*)^T b).$$

Let $\hat{\beta} = \operatorname{argmin}_b L_n(b)$ and let $\Delta_k = \operatorname{argmin}_{\Delta} L_{n,k}(\hat{\beta} + \Delta)$.

Proposition 1 (Heuristic claim). *If l is Lipschitz, smooth and symmetric and the data is “well-conditioned”, then $\Delta_k = O(1/n)$ (so (x_k, y_k) has a small influence on $\hat{\beta}$).*

Sketch of proof. Note that l being Lipschitz and differentiable implies that the derivatives of l are uniformly bounded by some constant C . Observe that

$$\begin{aligned}\nabla_{\Delta} L_{n,k}(\hat{\beta} + \Delta) &= \frac{1}{n} \sum_{i \neq k} l'(x_i^T(\hat{\beta} + \Delta) - y_i)x_i + \frac{1}{n} l'((x_k^*)^T(\hat{\beta} + \Delta) - y_k^*)x_k^* \\ &= \frac{1}{n} \sum_{i=1}^n l'(x_i^T(\hat{\beta} + \Delta) - y_i)x_i \\ &\quad + \frac{1}{n} \left(l'((x_k^*)^T(\hat{\beta} + \Delta) - y_k^*)x_k^* - l'(x_k^T(\hat{\beta} + \Delta) - y_k)x_k \right) \\ &= \nabla_{\Delta} L_n(\hat{\beta} + \Delta) + \frac{1}{n} (C_k^* x_k^* - C_k x_k),\end{aligned}$$

where $C_k^* = l'((x_k^*)^T(\hat{\beta} + \Delta) - y_k^*)$ and $C_k = l'(x_k^T(\hat{\beta} + \Delta) - y_k)$. We can do a Taylor's approximation of $\nabla_{\Delta} L_n(\hat{\beta} + \Delta)$ at $\Delta = 0$. We know that $\nabla_{\Delta} L_n(\hat{\beta} + \Delta)|_{\Delta=0} = 0$ since $\hat{\beta}$ minimizes $L_n(b)$. Thus

$$\nabla_{\Delta} L_{n,k}(\hat{\beta} + \Delta) \approx \nabla^2 L_n(\hat{\beta}) \Delta + \frac{1}{n} (C_k^* x_k^* - C_k x_k).$$

If we set the right hand side equal to 0 we get

$$\Delta_k \approx \nabla^2 L_n(\hat{\beta})^{-1} \frac{1}{n} (C_k^* x_k^* - C_k x_k).$$

This gives the approximation since l being Lipschitz implies that C_k^*, C_k are bounded and $\{x_i\}_{i=1}^n$ being well-conditioned implies that the inverse Hessian $\nabla^2 L_n(\hat{\beta})^{-1}$ exists and isn't too wild. \square

Definition 2. The *influence function* of an M-estimator is (roughly) the change in $\hat{\beta}^*$ when we change one observation. For M-estimators with $L(b) = \mathbb{E}[l(Y - X^T b)]$ and $\beta^* = \operatorname{argmin}_b L(b)$, we define the influence function ψ to be

$$\psi(x, y) = (\Delta^2 L(\beta^*))^{-1} l'(y - x^T \beta^*) x.$$

Theorem 1. Suppose $n \rightarrow \infty$. Let

$$\hat{\beta}(x, y) = \operatorname{argmin}_b \left\{ \sum_{i=1}^n l(y_i - x_i^T b) + l(y - x^T b) \right\},$$

and

$$\hat{\beta} = \operatorname{argmin}_b \left\{ \sum_{i=1}^n l(y_i - x_i^T b) \right\}.$$

Then, we have shown heuristically,

$$\hat{\beta}(x, y) = \hat{\beta} + \frac{1}{n} \psi(x, y) + o(1/n).$$


The upshot is that if the derivative $l'(t)$ is bounded, then the estimator is robust. If $l'(t)$ is unbounded, then the estimator is not robust. For example the following losses

$$\begin{aligned}l_1(t) &= \log(1 + e^t) + \log(1 + e^{-t}), \\ l_2(t) &= |t| \\ l_3(t) &= \begin{cases} \frac{1}{2} t^2 & \text{if } |t| \leq 1, \\ |t| - \frac{1}{2} & \text{if } |t| > 1. \end{cases}\end{aligned}$$

have corresponding derivatives which are all bounded on \mathbb{R}

$$\begin{aligned} l'_1(t) &= \frac{1}{1+e^{-t}} - \frac{1}{1+e^t}, \\ l'_2(t) &= \text{sgn}(t) \\ l'_3(t) &= \begin{cases} t & \text{if } |t| \leq 1, \\ \text{sgn}(t) & \text{if } |t| > 1. \end{cases} \end{aligned}$$

The derivatives look like this with l'_1 in black, l'_2 in blue and l'_3 in red. Note that l'_2 and l'_3 overlap for $|t| \geq 1$.



4 Quantile regression


Often, instead of predicting Y , we want an interval where Y is likely to be. That is, we would like a confidence set $\hat{c}(x) = [q_\alpha(x), q_{1-\alpha}(x)]$ such that $\mathbb{P}(Y \in \hat{c}(x) | X = x) = 1 - 2\alpha$ (this is impossible but let's try anyway). There are many applications where we are more interested in $\hat{c}(x)$ than \hat{y} . For example we might be

- Predicting election results,
- Predicting patient survival times,
- Predicting air quality.

What loss should we choose to get something like this? We could fit two (or more) models that predict the α , $1 - \alpha$ quantiles of $Y | X = x$. To do this we use the *pinball loss* (also called the quantile loss). For $\alpha \in (0, 1)$, let l_α be the loss function given by

$$l_\alpha(t) = \alpha(t)_+ + (1 - \alpha)(-t)_+,$$

where $(t)_+ = \max(0, t)$ so that $(t)_+ = 0$ for $t \leq 0$ and $(t)_+ = t$ for $t \geq 0$. The function l_α looks something like this. The blue loss is typical for $\alpha > 1/2$ and the black loss is typical for $\alpha < 1/2$. These losses are chosen with guiding principle (a) in mind. We will now show that the minimizers of $\mathbb{E}[l_\alpha(Y - t)]$ are α -quantiles of Y .



Define $R_\alpha(t) = \mathbb{E}[l_\alpha(Y - t)]$, then

$$R'_\alpha(t) = \mathbb{E}[-\alpha \mathbf{1}(Y > t) + (1 - \alpha) \mathbf{1}(Y \leq t)] = (1 - \alpha) \mathbb{P}(Y \leq t) - \alpha \mathbb{P}(Y > t).$$

If $\mathbb{P}(Y \leq t) > \alpha$, then $\mathbb{P}(Y > t) < 1 - \alpha$ and $R'_\alpha(t) > 0$ so t is too large. If $\mathbb{P}(Y \leq t) < \alpha$, then $\mathbb{P}(Y > t) > 1 - \alpha$ and $R'_\alpha(t) < 0$ so t is too small. Thus

$$\operatorname{argmin}_t R_\alpha(t) = \inf\{t : \mathbb{P}(Y \leq t) \geq \alpha\} = \alpha\text{-quantile of } Y.$$

Quantile regression is the following procedure. For $\alpha \in (0, 1/2)$, fit two models

$$\hat{\beta}_\alpha = \operatorname{argmin}_b \sum_{i=1}^n l_\alpha(y_i - x_i^T b) \quad \text{and} \quad \hat{\beta}_{1-\alpha} = \operatorname{argmin}_b \sum_{i=1}^n l_{1-\alpha}(y_i - x_i^T b).$$

Then our prediction intervals are

$$\hat{c}(x) = [\hat{\beta}_\alpha^T x, \hat{\beta}_{1-\alpha}^T x].$$

More generally, instead of linear functions we could consider fitting $f \in \mathcal{F}$ where \mathcal{F} is some function class. We would then have

$$\hat{f}_\alpha = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n l_\alpha(y_i - f(x_i)) \quad \text{and} \quad \hat{f}_{1-\alpha} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n l_{1-\alpha}(y_i - f(x_i)).$$

We would then again define

$$\hat{c}(x) = [\hat{f}_\alpha(x), \hat{f}_{1-\alpha}(x)].$$

Some issues:

- There is no guarantee that $\hat{f}_\alpha(x) \leq \hat{f}_{1-\alpha}(x)$. We can fix this by defining

$$\hat{c}(x) = [\min\{\hat{f}_\alpha(x), \hat{f}_{1-\alpha}(x)\}, \max\{\hat{f}_\alpha(x), \hat{f}_{1-\alpha}(x)\}].$$

- We do not have any guarantee that $\hat{c}(x)$ is a valid confidence interval for Y given $X = x$. That is we have no reason to believe that

$$\mathbb{P}(Y \in \hat{c}(x) | X = x) = 1 - \alpha.$$

The second issue cannot be resolved. It is a fact that without strong assumptions, you cannot have *conditional coverage*. That is we cannot find a procedure \hat{c} based on a i.i.d. sample $\{x_i, y_i\}_{i=1}^n$ such that

$$\mathbb{P}(Y_{n+1} \in \hat{c}(x_{n+1}) | X_{n+1} = x_{n+1}) = 1 - \alpha + o(1),$$

where Y_{n+1}, x_{n+1} is a new independent data point from the same distribution as our sample.

We can achieve *marginal coverage*. That is there exists a procedure \hat{c} such that

$$\mathbb{P}(Y_{n+1} \in \hat{c}(X_{n+1})) \geq 1 - \alpha.$$

Furthermore if Y is continuous, then we can find \hat{c} such that

$$\mathbb{P}(Y_{n+1} \in \hat{c}(X_{n+1})) = 1 - \alpha \pm \frac{1}{n}.$$

These ideas will be discussed on Thursday when we look at conformal confidence intervals.